

Abstract. Task-based parallel programming languages require the programmer to partition the traditional sequential code into smaller tasks in order to take advantage of the existing dataflow parallelism inherent in the applications. However, obtaining the partitioning that achieves optimal parallelism is not trivial because it depends on many parameters such as the underlying data dependencies and global problem partitioning. In order to help the process of finding a partitioning that achieves high parallelism, this paper introduces a framework that a programmer can use to: 1) estimate how much his application could benefit from dataflow parallelism; and 2) find the best strategy to expose dataflow parallelism in his application. Our framework automatically detects data dependencies among tasks in order to estimate the potential parallelism in the application. Furthermore, based on the framework, we develop an interactive approach to find the optimal partitioning of code. To illustrate this approach, we present a case study of porting High Performance Linpack from MPI to MPI/SMPs. The presented approach requires only superficial knowledge of the studied code and iteratively leads to the optimal partitioning strategy. Finally, the environment provides visualization of the simulated MPI/SMPs execution, thus allowing the developer to qualitatively inspect potential parallelization bottlenecks.

1 Introduction

New proposals for large-scale programming models are persistently spawning, but most of these initiatives fail because they attract little interest of the community. It takes a giant leap of faith for a programmer to take his already working sequential application and to port it to a novel programming model. This is problematic because the programmer cannot anticipate how would his application perform if it was ported to the new programming model, so he must decide whether the porting is worth the effort. Moreover, the programmer usually needs developing tools that would make the process of porting easier.

MPI/SMPs is a new hybrid dataflow programming model that should be efficient for numerous applications. In a manner similar to MPI, MPI/SMPs parallelizes computation of the distributed-memory nodes.

parallelism (parallelism of code sections that are mutually “far” from each other). Finally, MPI/SMPSSs outperforms MPI in numerous codes [8], among them the High Performance Linpack (HPL), the application that is used to rank parallel machines on the top 500 supercomputers lists [1].

To continue its progress, MPI/SMPSSs must get wider community in order to encourage MPI programmers to port their applications to MPI/SMPSSs. This encouragement is strictly related to assuring the programmer that he can achieve from this porting and that the porting would be easy. Therefore, our goal in this study is to develop a framework that provides support to:

- help an MPI programmer estimate how much parallelism MPI/SMPSSs can achieve in his MPI application, so he can decide whether the porting is worth the effort.
- help an MPI programmer find the optimal strategy to port his MPI application to MPI/SMPSSs.

2 SMPSSs Programming Model

SMPSSs [10] is a new shared-memory task-based parallel programming model that uses dataflow to exploit parallelism. SMPSSs slightly extends C, C++ and Fortran, offering semantics to declare some part of a code as a task. The programmer specifies memory regions on which that task operates. In porting a sequential code to SMPSSs, the programmer has to specify the following: *taskification* – to mark with pragma statements the functions that should be executed in parallel and *directionality of parameters* – to mark inside pragmas how are the parameters used within these functions. The specified directionality can be *input*, *output* and *inout*. Figure 1 illustrates the annotations needed to port a sequential C code to SMPSSs.

Given the annotations, the runtime is free to schedule all tasks out of order as long as the data dependencies are satisfied. The main thread starts execution. When it reaches a taskified function, it instantiates it as a task and proceeds with its own execution. Based on the parameters’ directionality, the runtime places

<pre>#pragma css task input(A[SizeA]) output (B[SizeB]) void compute(float *A, float *B) { ... }</pre>	<pre>int main () { ... compute(a,b); ... }</pre>
--	--

Note: The code in black presents the unchanged code of the legacy C application. The code in dark gray presents the annotations needed to mark the *taskification choice*, while the code in light gray presents the annotations needed to declare the *directionality of parameters*.

Fig. 1. Annotations needed to port a code from sequential C to SMPSSs.

automatically renames data objects to avoid all false dependencies (dependencies caused by buffer reuse).

Integrated with MPI, SMPSSs allows to taskify functions with MPI and thus potentially extract very distant parallelism. The idea is to embed functions with MPI transfers inside tasks, and thus relate the messaging to dataflow dependencies. For example, a task with *MPI_Send* of some buffer locally reads (*input* directionality) that buffer from the memory and transfers it to the network, while a task with *MPI_Recv* of some buffer gets the data from the network and locally stores (*output* directionality) it to the memory. Taskification of transfers overcomes strong synchronization points of sequential execution and potentially exploits distant parallelisms, providing much more flexible messaging behavior than fork-join based MPI/OpenMP. Marjanovic et al. showed that apart from better peak GFlops/s performance, compared to sequential MPI/SMPSSs delivers better tolerance to bandwidth reduction and external perturbations (such as OS noise).

3 Motivation

Finding the best taskification strategy is far from trivial. Figure 2 shows a sequential application composed of four computational parts (*A*, *B*, *C*, *D*), the data dependencies among those parts, and some of the possible taskification strategies. Although the application is very simple, it allows many possible taskifications that expose different amount of parallelism. *T0* puts all code in a single task and, in fact, presents non-SMPSSs code. *T1* and *T2* both break the application into two tasks but fail to expose any parallelism. On the other hand, *T3* and *T4* both break the application into 3 tasks, but while *T3* achieves no parallelism, *T4* exposes parallelism between *C* and *D*. Finally, *T5* breaks the application into 4 tasks but achieves the same amount of parallelism as *T4*. Considering that increasing the number of tasks increases the runtime overhead of creating and scheduling tasks, one can conclude that the optimal taskification is *T4*, because it gives the highest speedup with the lowest cost of the number of tasks. On the other hand, for a complex MPI application, the number of possible taskifications could be huge, so finding the optimal taskification is a non-trivial task.




Fig. 2. Execution of different possible taskifications for a code composed of

the input code in the pure MPI code with inserted functions annotations and exits from tasks. Then the obtained code is compiled and executed in MPI fashion. Each MPI process runs on top of one instance of Valgrind machine that implements a designed tracer. The tracer makes the trace of (actually executed) MPI execution, while at the same time, it reconstructs would be the traces of the (potential) MPI/SMPSSs execution. Dimemator merges the obtained traces and reconstructs time-behavior of these on a parallel platform. Finally, Paraver can visualize the simulated time-behavior and allow to profoundly study the differences between the (instrumented) and the (corresponding simulated) MPI/SMPSSs execution. In our previous work [14], we used a similar idea to estimate the potential benefits of overlapped communication and computation in pure MPI applications.

4.1 Input Code

The input code can be MPI/SMPSSs code or an MPI code with light annotations. The input code has to specify which functions (parts of code) should be executed as tasks, but not the directionality of the function parameters. Thus, the input code can be an MPI code, only with annotations specifying which parts of code should be executed as tasks. Figure 4 on the left shows an example of input code with annotated *taskification* choice.

4.2 Code Translator

Our Mercurium based tool translates the input code into the code with task serialization of tasks. The obtained code is a pure MPI code with empty annotations (*hooks*) annotating when the execution enters and exits from a task (hook). The translated code is then compiled with *mpicc*, and the binary of the code execution is passed for further instrumentation. It is important to note

Input code	Translated code
<pre>#pragma css task void compute(float *A, float *B) { ... } int main () { ... compute(a,b); ... }</pre>	<pre>void compute(float *A, float *B) { ... } int main () { ... start_task_valgrind("compute"); compute(a,b); end_task_valgrind("compute"); ... }</pre>

Note: The input code does not have to be a complete MPI/SMPSSs code, because the input code only needs to mark all entries/exits from each task. Thus, as shown, the input code can be an MPI application only with a specified proposed taskification.

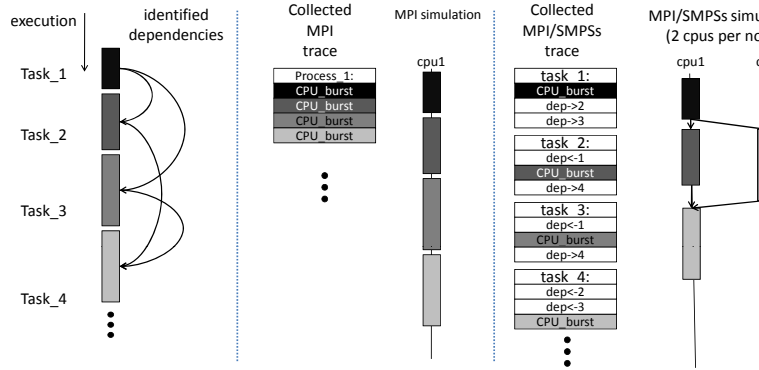
Fig. 4. Translation of the input code required by the framework

4.3 Tracer

Valgrind [9] is a virtual machine that uses just-in-time (JIT) compilation techniques. The original code of an application never runs directly on the processor. Instead, the code is first translated into a temporary, simpler, portable neutral form called Intermediate Representation (IR). Then, the developer is free to do any translation of the IR, before Valgrind translates the IR into machine code and lets the host processor run it.

Leveraging Valgrind functionalities, the tracer instruments the execution of the application. It makes two Dimemas traces: one describing the instrumented MPI execution and the other describing the potential MPI/SMPSSs execution. The tracer uses the following Valgrind functionalities: 1) intercepting the inserted hooks in order to track which task is currently being executed; 2) intercepting all dynamic allocations in order to maintain the pool of data objects in the memory; 3) intercepting memory accesses in order to identify data dependencies among tasks; and 4) intercepting all MPI calls in order to track MPI activity of the execution. Using the obtained information, the tracer generates the trace of the (actually executed) MPI execution, while at the same time, it reconstructs what would be the trace of the potential (not executed) MPI/SMPSSs execution.

The tool instruments accesses to all memory objects and derives data dependencies among tasks. By intercepting all dynamic allocations and releases of memory (*allocs* and *frees*), the tool maintains the pool of all dynamic memory objects. Similarly, by intercepting all static allocations and releases



Note: The tracer describes the MPI traces by emitting two types of records: 1) computation records specifying the length of computation burst; and 2) communication record specifying the parameters of MPI transfers. Conversely, it describes the MPI/SMPSSs trace by breaking the original computation bursts into tasks and synchronizing the created tasks according to the identified data dependencies.

Fig. 5. Collecting trace of the original MPI and the potential MPI/SMPSSs

at the granularity of one byte. Based on these records, and knowing the task the execution is at every moment, the tracer detects all read-after-write dependencies and interpret them as dependencies among tasks.

The tool creates the trace of the executed MPI run, and at the same time, considering identified task dependencies, it creates what would be the potential MPI/SMPSSs run (Figure 5). When generating the original MPI/SMPSSs run, the tool describes the actually executed run by putting in the trace two records: 1) computation record stating the length of computation bursts and 2) the number of instructions 2) communication record specifying the parameters of the executed MPI transfer. Additionally, when reconstructing the potential MPI/SMPSSs run, the tracer breaks the original computation into tasks, and then synchronizes the created tasks according to the data dependencies.

4.4 Replay Simulator

Dimemas is an open-source tracefile-based simulator for analysis of passing applications on a configurable parallel platform. The communication model, validated in [4], consists of a linear model and nonlinear effects such as network congestion. The interconnect is parametrized by bandwidth and the number of global buses (denoting how many messages can communicate travel throughout the network). Also, each processor is characterized by the number of input/output ports that determine its injection rate to the network. Finally, the simulated output of Dimemas can be visualized in Paraver.

We extended Dimemas to support synchronization of tasks in a way that allows Paraver to visualize all data dependencies. We implemented a synchronization using an intra-node instantaneous MPI transfer that synchronizes the source and the destination tasks. This way, Paraver can visualize the execution time-behavior showing both MPI communications among processes and data dependencies among tasks. Using this feature, the developer can visualize each execution bottleneck and further inspect its causes.

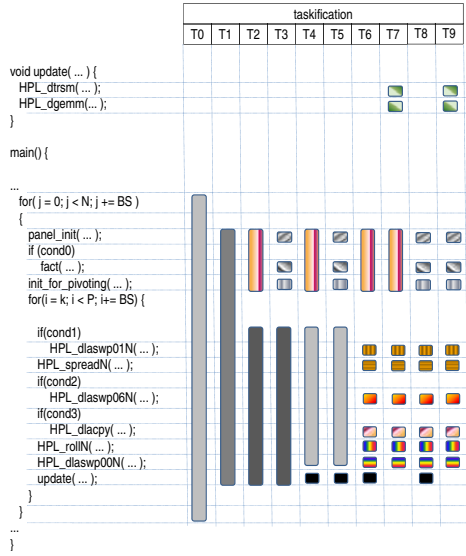
5 Experiments

Our experiments explore MPI/SMPSSs execution of HP Linpack on a many-core nodes. We used HPL with the problem size of 8192 and (PxQ) data decomposition. Also, we test various granularities of execution running HPL with block sizes (BS) of 32, 64, 128, and 256. Our target machine consists of four many-core nodes, with one MPI process running on each node. We are primarily interested in the MPI/SMPSSs potential parallelism in the code, so we make most of the measurements for unlimited resources. The target machine – infinite number of cores per node and ideal interconnect.

The major part of our experiment consist of exploring the potential SMPs taskifications of HPL. In a case study with HPL, we present a bottom-to-top approach that uses a trial-and-error method, requires no knowledge of the studied code, and finally leads to exposing dataflow parallelism in the code. The approach uses the following method: 1) we propose a coarse-grained taskification for the code; 2) given the taskification, the environment determines the potential speedup and offers visualization of the resulting MPI/SMP taskification. 3) based on the output, we choose a finer-grained taskification and go to step 2. We start from the most coarse-grain taskification (T_0) that presents one MPI process into one task and actually presents the traditional MPI taskification (Figure 6(a)). Then using T_0 as the baseline, we determine the *potential speedup* of T_i ($1 \leq i \leq 9$) *normalized to T_0* as the speedup of T_i over T_0 when both these taskifications execute on a machine with unlimited number of processors per node and unlimited network performance (Figure 7(b)).

5.1 Results

First, the framework instruments the application to obtain the profile of the code during the taskification process. Table 6(b) shows the accumulated time spent in tasks.



(a) HPL and the evaluated taskifications.

			granul	
			BS-32	BS-64
task name	outer	panel_init	0.0003	0.0002
		fact	0.7525	1.2071
		init_for_pivoting	0.0246	0.0487
	inner	HPL_dlaswp01N	0.2583	0.2917
		HPL_sreadN	0.1599	0.0800
		HPL_dlaswp06N	0.1222	0.1359
		HPL_rolN	0.3267	0.1619
		HPL_dlacpy	0.0857	0.0932
		HPL_dlaswp00N	0.3706	0.4485
	update	HPL_dtrsm	0.8269	1.6674
		HPL_dgemm	97.0683	95.8614

(b) Distribution of total time spent in tasks (%).

			granul	
			BS-32	BS-64
task name	outer	panel_init	0.0003	0.0003
		fact	1.3468	3.7670
		init_for_pivoting	0.0221	0.0761
	inner	HPL_dlaswp01N	0.0073	0.0289
		HPL_sreadN	0.0022	0.0040
		HPL_dlaswp06N	0.0034	0.0135
		HPL_rolN	0.0046	0.0080
		HPL_dlacpy	0.0024	0.0092
	update	HPL_dlaswp00N	0.0052	0.0222
		HPL_dtrsm	0.0117	0.0826
		HPL_dgemm	1.3677	4.7492

(c) Average function duration.

Note: In Tables 6(b) and 6(c), apart from statistic for each function of the code, we provide statistics for two logical sections: **outer** – consisting of `panel_init`, `fact` and `init_for_pivoting`; **inner** – consisting of `HPL_dlaswp01N`, `HPL_sreadN`, `HPL_dlaswp06N`, `HPL_rolN`, `HPL_dlacpy`, `HPL_dlaswp00N`, `HPL_dtrsm` and `HPL_dgemm`.


Fig. 6. Taskifications evaluated for HPL and duration and time spent in each task.

spends in that function from 95.57% (for $BS = 256$) to 97.83% (for $BS = 128$). On the other hand, Figure 6(c) shows the average duration of each function. This information identifies which function is a good candidate to be broken down into smaller tasks. In this example, function *panel_init* is very short so breaking it into smaller tasks makes little sense. Also, it is important to note that *BS* reduces execution time of most of the functions, so this could also be used to make finer-grained execution.

Considering the data showed on previous tables, we start the process of proposing parallelism by: 1) proposing a taskification ($T1 - T9$ in Figure 7(a)); 2) testing how many tasks we created (Figure 7(a)); and 3) testing the speedup of the taskification (Figure 7(b)). $T0$ is the baseline taskification where each iteration makes only one task per MPI process. $T1$ puts each iteration of the outer loop in one task, but this strategy gives no additional parallelism compared to $T0$. Furthermore, $T2$ breaks down the code into section *outer* and separate the inner loop, still giving no improvement in speedup. $T3$ additionally breaks down the inner loop, still giving no improvement in speedup.

		granularity			
		BS-32	BS-64	BS-128	BS-256
taskification	T1	1024	516	260	132
	T2	66.314	16.778	4.298	1.130
	T3	69.898	18.570	5.194	1.578
	T4	131.600	33.040	8.336	2.128
	T5	135.184	34.832	9.232	2.576
	T6	327.458	81.826	20.450	5.122
	T7	425.382	106.216	26.502	6.614
	T8	331.042	83.618	21.346	5.570
	T9	428.966	108.006	27.398	7.062

(a) Total number of tasks created.



(b) Speedup normalized to $T0$.

Note: In Figure 7(b), all taskifications ($T0-T9$) execute in MPI/SMPs fashion on an 8-core machine. Then, the speedup of taskification Ti over taskification $T0$ represents the parallelism of taskification Ti normalized to taskification $T0$.

Fig. 7. Number of task instances and the potential parallelism of each taskification.





Fig. 8. Paraver visualization of the first 63 tasks and the dependencies among them (taskification $T4$, $BS=256$)


thus impedes parallelism. Finally, $T4$ compared to $T2$ separates sections $outer$ and $inner$ into two parallel sections, thus releasing the significant amount of parallelism. Finally, $T4$ compared to $T2$ separates sections $outer$ and $inner$ into two parallel sections, thus releasing the significant amount of parallelism. it achieves the speedup of 6.76, 12.28, 21.48 and 32.02 for block size of 256, 128, 64 and 32, respectively (Figure 7(b)). Also, $T4$ significantly increases the number of tasks in the application to 2.128, 8.336, 33.040 and 131.600 for block sizes of 256, 128, 64 and 32, respectively (Figure 7(a)). Now, Paravision reveals that in $T4$: 1) each section $inner$ depends on the section $outer$ in the previous iteration of the inner loop; and 2) each $update$ depends on the section $inner$ in the same iteration of the inner loop. Thus, because section $inner$ is much shorter than $update$, all dependent sections $inner$ can execute concurrently and then independent instances of $update$ can execute concurrently (Figure 7(b)).

Further breaking down of $outer$, $inner$ and $update$ contributes little to the potential speedup (Figure 7(b)). Breaking of $outer$, for block sizes of 256 and 64, causes slightly higher parallelism of $T5$, $T8$ and $T9$, compared to $T4$, $T6$ and $T7$. On the other hand, breaking of $inner$, for block size of 32, causes significantly higher parallelism of $T6$, $T7$, $T8$ and $T9$, compared to $T4$, $T5$, $T7$ and $T8$. This effect happens because for very high concurrency of $update$ (speedup more than 30), the critical path of the execution moves and starts passing through section $inner$. In these circumstances breaking of $inner$ significantly increases parallelism by allowing concurrency of functions $HPL_dlaswp00N$, $HPL_dlaswp01N$ and $HPL_dlaswp06N$. Finally, breaking of $update$, for block size 32, causes higher parallelism of $T9$ compared to $T8$.

Figure 9 shows the speedup and parallel efficiency of $T9$ for different numbers of cores per node. The results show that high parallelism in the application is not useful not to achieve high speedup on a small parallel machine, but to deploy efficiently a large parallel machine. Figure 9(a) shows that for a



(a) Speedup.



(b) Parallel efficiency.

Note: Parallel efficiency denotes the ratio between the application's speedup achieved on a parallel machine and the number of cores of that parallel machine. Infact, the metric is the overall average core utilization in the whole machine.

Fig. 9. Speedup and parallel efficiency for $T9$ for various number of cores per node.

achieves the speedup of 6.80, 12.34, 21.57 and 29.47, respectively. Further, Figure 9(b) shows parallel efficiency (core utilization) – the ratio between application’s speedup achieved on some parallel machine and the number of cores in that machine. Adopting that an application efficiently utilizes a machine if the parallel efficiency is higher than 75%, the results show that T9 with problem sizes of 256, 128, 64 and 32, can efficiently utilize the machine of 8, 15, 24 and 32 cores per node, respectively. Therefore, to efficiently employ many-core machines with hundreds of cores per node, HPL has to expose even more parallelism. For instance, by making finer-grain taskification with further reduction of block size.

6 Related Work

Back in 1991 the community started claiming that instruction-level parallelism is dead [15], and consequently in the following 20 years appeared many programming models that exploit task-level parallelism. OpenMP [11] is a popular programming model for shared memory that was founded with the idea of parallelizing loops, but from version 3.0 provides support for task parallelism. Cilk [2] implements a model of spawning various tasks and specifying synchronization point where these tasks are waited for. MPI tasklets [5] parallelizes SMP tasks by incorporating dynamic scheduling strategy into current MPI implementations. There are also proposals that originated from the industry, such as: TBB [12] from Intel and TPL [6] from Microsoft. Still, all these models suffer from the limitations of fork-join based programming models. On the other hand, SMPs [10] is a programming model in which the programmer specifies dependencies among tasks, rather than specifying synchronization points. Based on the specified dependencies, the runtime schedules tasks in a dataflow manner, potentially extracting very distant parallelism. Furthermore, SMPs can be integrated with MPI, allowing better messaging behavior. M. J. Frumkin *et. al.* [8] demonstrate that compared to MPI, MPI/SMPs provides better performance as well as higher tolerance to network reduction and external interference.

However, there is little development support for these programming models. Alchemist tool [16] identifies parts of code that are suitable for task-level speculation. Embla [7] estimates the potential speed-up of fork-join based parallelization. Starsscheck [3] checks correctness of pragma annotations for a family of programming models. Our work adds up to these efforts by developing a framework that estimates the potential parallelism of MPI/SMPs. Furthermore, our work goes beyond the state-of-the-art tools because: 1) it deals with a new execution model that integrates MPI with task-based dataflow execution; 2) it allows to study MPI/SMPs execution before the original MPI application is ported to MPI/SMPs; 3) it provides an estimation of the parallelism on a configurable target platform; and 4) it provides visualization of the task execution.

tional parallelism inherent in MPI parallel programs. However, the cost of this type of execution impedes an MPI programmer from anticipating the dataflow parallelism he can obtain in his application. Moreover, it is not clear how to determine which parts of code should be encapsulated into tasks in order to expose the parallelism and still avoid creating unnecessary tasks that increase runtime overhead. To address this issue, we have developed a framework that automatically estimates the potential dataflow parallelization in applications. We show how, using the framework, one can find optimal taskification for any application through a trial-and-error iterative approach that requires no prior knowledge of the studied code. We prove the effectiveness of this approach in a case study in which we explore the taskification of High Performance Linpack (HPL). The results show that HPL expresses substantial amount of dataflow parallelism that allows the application to efficiently utilize all available nodes with up to 47 cores per node. Moreover, we show that the granularity of tasking significantly impacts parallel efficiency, and thus, in order to efficiently utilize higher number of cores, finer-granularity of execution should be used.

Acknowledgements. We thankfully acknowledge the support of the European Commission through the HiPEAC-2 Network of Excellence (FP7/ICT-2007-215869), the TEXT project (IST-2007-261580), and the support of the Spanish Ministry of Education (TIN2007-60625, and CSD2007-00050), and the Generalitat de Catalunya (2009-SGR-980).

References

1. Top500 List: List of top 500 supercomputers, <http://www.top500.org/>
2. Blumofe, R.D., Joerg, C.F., Kuszmaul, B.C., Leiserson, C.E., Randall, K.Y.: Cilk: An Efficient Multithreaded Runtime System. *J. Parallel Distributed Comput.* 37, 55–69 (1996)
3. Carpenter, P.M., Ramirez, A., Ayguade, E.: Starsscheck: A tool to find task-based parallel programs. In: D’Ambra, P., Guarracino, M., Talia, D. (eds.) *Euro-Par 2010*. LNCS, vol. 6271, pp. 2–13. Springer, Heidelberg (2010)
4. Girona, S., Labarta, J., Badia, R.M.: Validation of dimemas communication for mpi collective operations. In: *PVM/MPI*, pp. 39–46 (2000)
5. Kale, V., Gropp, W.: Load Balancing for Regular Meshes on SMPs with MPI. In: Keller, R., Gabriel, E., Resch, M., Dongarra, J. (eds.) *EuroMPI 2007*. LNCS, vol. 6305, pp. 229–238. Springer, Heidelberg (2010)
6. Leijen, D., Hall, J.: Parallel performance: Optimize managed code for multi-processor machines. *MSDN Magazine* (2007)
7. Mak, J., Faxén, K.-F., Janson, S., Mycroft, A.: Estimating and Exploiting Dataflow Parallelism by Source-Level Dependence Profiling. In: D’Ambra, P., Guarracino, M., Talia, D. (eds.) *Euro-Par 2010*. LNCS, vol. 6271, pp. 26–37. Springer, Heidelberg (2010)

10. Perez, J.M., Badia, R.M., Labarta, J.: A dependency-aware task-based programming environment for multi-core architectures. In: CLUSTER, pp. 142–153 (2008)
11. Proposed Industry Standard. Openmp: A proposed industry standard for shared memory programming
12. Reinders, J.: Intel threading building blocks: outfitting C++ for multi-processor parallelism. O'Reilly Media, Inc., Sebastopol (2007)
13. Snir, M., Otto, S., Huss-Lederman, S., Walker, D., Dongarra, J.: MPI: A Complete Reference. The MIT Press, Cambridge (1998)
14. Subotic, V., Sancho, J.C., Labarta, J., Valero, M.: A Simulation Framework to Automatically Analyze the Communication-Computation Overlap in Scientific Applications. In: CLUSTER 2010 (2010)
15. Wall, D.W.: Limits of Instruction-Level Parallelism. In: ASPLOS (1991)
16. Zhang, X., Navabi, A., Jagannathan, S.: Alchemist: A transparent distributed distance profiling infrastructure. In: CGO 2009 (2009)